

# Week7\_SimonsenHomework

Steven Simonsen

2024-10-11

```
library(dplyr)

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(caret)

## Loading required package: ggplot2
## Loading required package: lattice

library(reticulate)
library(tensorflow)

##
## Attaching package: 'tensorflow'
## The following object is masked from 'package:caret':
##
##   train

library(keras3)

##
## Attaching package: 'keras3'
## The following objects are masked from 'package:tensorflow':
##
##   set_random_seed, shape

library(MESS)

setwd("C:\\Users\\steve\\OneDrive\\Documents\\School\\DSE6211\\Week7")

data <- read.csv("lab_7_data.csv")

training_ind <- createDataPartition(data$lodgepole_pine,
                                     p = 0.75,
                                     list = FALSE,
                                     times = 1)
```

```

training_set <- data[training_ind, ]
test_set <- data[-training_ind, ]

top_20_soil_types <- training_set %>%
  group_by(soil_type) %>%
  summarise(count = n()) %>%
  arrange(desc(count)) %>%
  select(soil_type) %>%
  top_n(20)

## Selecting by soil_type
training_set$soil_type <- ifelse(training_set$soil_type %in%
                                top_20_soil_types$soil_type,
                                training_set$soil_type,
                                "other")

training_set$wilderness_area <- factor(training_set$wilderness_area)
training_set$soil_type <- factor(training_set$soil_type)

onehot_encoder <- dummyVars(~ wilderness_area + soil_type,
                           training_set[, c("wilderness_area",
                                             "soil_type")],
                           levelsOnly = TRUE,
                           fullRank = TRUE)

onehot_enc_training <- predict(onehot_encoder,
                              training_set[, c("wilderness_area",
                                                "soil_type")])

training_set <- cbind(training_set, onehot_enc_training)

test_set$soil_type <- ifelse(test_set$soil_type %in%
                             top_20_soil_types$soil_type,
                             test_set$soil_type,
                             "other")

test_set$wilderness_area <- factor(test_set$wilderness_area)
test_set$soil_type <- factor(test_set$soil_type)

onehot_enc_test <- predict(onehot_encoder, test_set[,
                                                    c("wilderness_area",
                                                      "soil_type")])

test_set <- cbind(test_set, onehot_enc_test)

test_set[, -c(11:13)] <- scale(test_set[, -c(11:13)],
                              center = apply(training_set[,
                                                    -c(11:13)],
                                              2, mean),
                              scale = apply(training_set[,
                                                    -c(11:13)],
                                              2, sd))

```

```

training_set[, -c(11:13)] <- scale(training_set[, -c(11:13)])

training_features <- array(data = unlist(training_set[,
                                          -c(11:13)]),
                           dim = c(nrow(training_set), 33))

training_labels <- array(data = unlist(training_set[, 13]),
                          dim = c(nrow(training_set)))

test_features <- array(data = unlist(test_set[, -c(11:13)]),
                       dim = c(nrow(test_set), 33))

test_labels <- array(data = unlist(test_set[, 13]),
                     dim = c(nrow(test_set)))


pca_results <- prcomp(training_features[, 1:10])
summary(pca_results)

## Importance of components:
##              PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation    1.6298 1.4433 1.3172 1.0353 0.87873 0.73984 0.68168
## Proportion of Variance 0.2656 0.2083 0.1735 0.1072 0.07722 0.05474 0.04647
## Cumulative Proportion 0.2656 0.4739 0.6475 0.7547 0.83186 0.88660 0.93307
##              PC8      PC9      PC10
## Standard deviation    0.60035 0.5532 0.05351
## Proportion of Variance 0.03604 0.0306 0.00029
## Cumulative Proportion 0.96911 0.9997 1.00000

screeplot(pca_results, type = "line")

```

## pca\_results



```
training_rotated <- as.matrix(training_features[, 1:10]) %*%  
  pca_results$rotation  
  
training_features <- cbind(training_features, training_rotated[,  
  1:6])  
  
test_rotated <- as.matrix(test_features[, 1:10]) %*%  
  pca_results$rotation  
  
test_features <- cbind(test_features, test_rotated[, 1:6])  
  
use_virtualenv("my_tf_workspace")  
  
model <- keras_model_sequential() %>%  
  layer_dense(units = 50, activation = "relu") %>%  
  layer_dense(units = 25, activation = "relu") %>%  
  layer_dense(units = 1, activation = "sigmoid")  
  
compile(model,  
  optimizer = "rmsprop",  
  loss = "binary_crossentropy",  
  metrics = "accuracy")  
  
history <- fit(model, training_features, training_labels,
```

```
epochs = 40, batch_size = 512,  
validation_split = 0.33)
```

```
## Epoch 1/40  
## 9/9 - 1s - 139ms/step - accuracy: 0.5832 - loss: 0.6810 - val_accuracy: 0.4810 - val_loss: 0.7419  
## Epoch 2/40  
## 9/9 - 0s - 18ms/step - accuracy: 0.6894 - loss: 0.6003 - val_accuracy: 0.5185 - val_loss: 0.7187  
## Epoch 3/40  
## 9/9 - 0s - 9ms/step - accuracy: 0.7239 - loss: 0.5625 - val_accuracy: 0.5385 - val_loss: 0.7196  
## Epoch 4/40  
## 9/9 - 0s - 10ms/step - accuracy: 0.7518 - loss: 0.5353 - val_accuracy: 0.5579 - val_loss: 0.7139  
## Epoch 5/40  
## 9/9 - 0s - 21ms/step - accuracy: 0.7643 - loss: 0.5145 - val_accuracy: 0.5857 - val_loss: 0.7046  
## Epoch 6/40  
## 9/9 - 0s - 8ms/step - accuracy: 0.7721 - loss: 0.4971 - val_accuracy: 0.5918 - val_loss: 0.7070  
## Epoch 7/40  
## 9/9 - 0s - 20ms/step - accuracy: 0.7776 - loss: 0.4834 - val_accuracy: 0.6200 - val_loss: 0.6859  
## Epoch 8/40  
## 9/9 - 0s - 7ms/step - accuracy: 0.7879 - loss: 0.4722 - val_accuracy: 0.6223 - val_loss: 0.6915  
## Epoch 9/40  
## 9/9 - 0s - 8ms/step - accuracy: 0.7888 - loss: 0.4636 - val_accuracy: 0.6237 - val_loss: 0.6982  
## Epoch 10/40  
## 9/9 - 0s - 8ms/step - accuracy: 0.7885 - loss: 0.4570 - val_accuracy: 0.6469 - val_loss: 0.6674  
## Epoch 11/40  
## 9/9 - 0s - 10ms/step - accuracy: 0.7942 - loss: 0.4509 - val_accuracy: 0.6274 - val_loss: 0.6984  
## Epoch 12/40  
## 9/9 - 0s - 8ms/step - accuracy: 0.7945 - loss: 0.4468 - val_accuracy: 0.6492 - val_loss: 0.6814  
## Epoch 13/40  
## 9/9 - 0s - 10ms/step - accuracy: 0.7952 - loss: 0.4435 - val_accuracy: 0.6358 - val_loss: 0.7005  
## Epoch 14/40  
## 9/9 - 0s - 25ms/step - accuracy: 0.8027 - loss: 0.4391 - val_accuracy: 0.6450 - val_loss: 0.6889  
## Epoch 15/40  
## 9/9 - 0s - 9ms/step - accuracy: 0.8029 - loss: 0.4355 - val_accuracy: 0.6511 - val_loss: 0.6862  
## Epoch 16/40  
## 9/9 - 0s - 8ms/step - accuracy: 0.8038 - loss: 0.4334 - val_accuracy: 0.6520 - val_loss: 0.6941  
## Epoch 17/40  
## 9/9 - 0s - 7ms/step - accuracy: 0.8059 - loss: 0.4305 - val_accuracy: 0.6525 - val_loss: 0.6869  
## Epoch 18/40  
## 9/9 - 0s - 8ms/step - accuracy: 0.8054 - loss: 0.4277 - val_accuracy: 0.6409 - val_loss: 0.7050  
## Epoch 19/40  
## 9/9 - 0s - 8ms/step - accuracy: 0.8091 - loss: 0.4255 - val_accuracy: 0.6497 - val_loss: 0.6950  
## Epoch 20/40  
## 9/9 - 0s - 9ms/step - accuracy: 0.8095 - loss: 0.4225 - val_accuracy: 0.6399 - val_loss: 0.7045  
## Epoch 21/40  
## 9/9 - 0s - 10ms/step - accuracy: 0.8116 - loss: 0.4211 - val_accuracy: 0.6367 - val_loss: 0.7182  
## Epoch 22/40  
## 9/9 - 0s - 8ms/step - accuracy: 0.8150 - loss: 0.4189 - val_accuracy: 0.6399 - val_loss: 0.7085  
## Epoch 23/40  
## 9/9 - 0s - 10ms/step - accuracy: 0.8150 - loss: 0.4172 - val_accuracy: 0.6325 - val_loss: 0.7157  
## Epoch 24/40  
## 9/9 - 0s - 11ms/step - accuracy: 0.8155 - loss: 0.4145 - val_accuracy: 0.6381 - val_loss: 0.7023  
## Epoch 25/40  
## 9/9 - 0s - 8ms/step - accuracy: 0.8141 - loss: 0.4139 - val_accuracy: 0.6358 - val_loss: 0.7053  
## Epoch 26/40
```

```

## 9/9 - 0s - 21ms/step - accuracy: 0.8169 - loss: 0.4113 - val_accuracy: 0.6214 - val_loss: 0.7324
## Epoch 27/40
## 9/9 - 0s - 11ms/step - accuracy: 0.8155 - loss: 0.4105 - val_accuracy: 0.6311 - val_loss: 0.7254
## Epoch 28/40
## 9/9 - 0s - 18ms/step - accuracy: 0.8196 - loss: 0.4088 - val_accuracy: 0.6247 - val_loss: 0.7279
## Epoch 29/40
## 9/9 - 0s - 7ms/step - accuracy: 0.8164 - loss: 0.4065 - val_accuracy: 0.6487 - val_loss: 0.6991
## Epoch 30/40
## 9/9 - 0s - 9ms/step - accuracy: 0.8191 - loss: 0.4059 - val_accuracy: 0.6409 - val_loss: 0.7151
## Epoch 31/40
## 9/9 - 0s - 7ms/step - accuracy: 0.8155 - loss: 0.4058 - val_accuracy: 0.6293 - val_loss: 0.7287
## Epoch 32/40
## 9/9 - 0s - 10ms/step - accuracy: 0.8226 - loss: 0.4033 - val_accuracy: 0.6348 - val_loss: 0.7215
## Epoch 33/40
## 9/9 - 0s - 9ms/step - accuracy: 0.8212 - loss: 0.4016 - val_accuracy: 0.6242 - val_loss: 0.7444
## Epoch 34/40
## 9/9 - 0s - 10ms/step - accuracy: 0.8242 - loss: 0.4004 - val_accuracy: 0.6247 - val_loss: 0.7406
## Epoch 35/40
## 9/9 - 0s - 10ms/step - accuracy: 0.8244 - loss: 0.3993 - val_accuracy: 0.6251 - val_loss: 0.7360
## Epoch 36/40
## 9/9 - 0s - 10ms/step - accuracy: 0.8262 - loss: 0.3976 - val_accuracy: 0.6270 - val_loss: 0.7353
## Epoch 37/40
## 9/9 - 0s - 9ms/step - accuracy: 0.8244 - loss: 0.3960 - val_accuracy: 0.6344 - val_loss: 0.7334
## Epoch 38/40
## 9/9 - 0s - 9ms/step - accuracy: 0.8274 - loss: 0.3958 - val_accuracy: 0.6348 - val_loss: 0.7316
## Epoch 39/40
## 9/9 - 0s - 9ms/step - accuracy: 0.8253 - loss: 0.3943 - val_accuracy: 0.6177 - val_loss: 0.7745
## Epoch 40/40
## 9/9 - 0s - 8ms/step - accuracy: 0.8258 - loss: 0.3932 - val_accuracy: 0.6330 - val_loss: 0.7406

```

```

predictions <- predict(model, test_features)

```

```

## 69/69 - 0s - 2ms/step

```

```

test_set$p_prob <- predictions[, 1]

```

```

##### ROC curve

```

```

roc_data <- data.frame(threshold=seq(1,0,-0.01), fpr=0, tpr=0)

```

```

for (i in roc_data$threshold) {
  over_threshold <- test_set[test_set$p_prob >= i, ]
  fpr <- sum(over_threshold$lodgepole_pine==0)/
    sum(test_set$lodgepole_pine==0)
  roc_data[roc_data$threshold==i, "fpr"] <- fpr
  tpr <- sum(over_threshold$lodgepole_pine==1)/
    sum(test_set$lodgepole_pine==1)
  roc_data[roc_data$threshold==i, "tpr"] <- tpr
}

```

```

ggplot() +
  geom_line(data = roc_data, aes(x = fpr, y = tpr, color =
    threshold), size = 2) +
  scale_color_gradientn(colors = rainbow(3)) +
  geom_abline(intercept = 0, slope = 1, lty = 2) +
  geom_point(data = roc_data[seq(1, 101, 10), ], aes(x = fpr, y =

```

```


tpr)) +
  geom_text(data = roc_data[seq(1, 101, 10), ],
    aes(x = fpr, y = tpr, label = threshold, hjust = 1.2,
      vjust = -0.2))

```

```

## Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
## i Please use `linewidth` instead.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.

```



```

##### AUC
auc <- auc(x = roc_data$fpr, y = roc_data$tpr, type = "spline")

## Warning in regularize.values(x, y, ties, missing(ties)): collapsing to unique
## 'x' values
auc

## [1] 0.8352148

```

## Exercises

Hint for all exercises: see the synchronous live session slides.

1) What is the main difference between unsupervised and supervised learning?

In supervised learning, each observation is associated with a label. In unsupervised learning, there are no

labels associated with the observations.

2) Is centering required for PCA? Is scaling required for PCA? Explain your answers.

Centering is always required for PCA since PCs are unit vectors originating from the origin. Scaling is almost always required to avoid PCA finding the features with the most variance. The exception is if numerical features are measured on similar scales and we want features with larger variance to have more importance.

3) After running the code above, run the following code to obtain the PCA feature loadings. Copy and paste the output. Which feature has the strongest influence on the first PC? Which feature has the weakest influence on the first PC?

```
pca_results$rotation
```

```
##          PC1          PC2          PC3          PC4          PC5          PC6
## [1,]  0.13408619  0.35811179 -0.33839272  0.033709165 -0.63116833 -0.53751583
## [2,]  0.48351213 -0.16505835  0.08552072  0.079881768 -0.03061640  0.03863554
## [3,] -0.11363833 -0.52197760 -0.15333845  0.378800022 -0.16016438  0.03699555
## [4,]  0.11426905  0.02682936 -0.64515588 -0.231499324  0.16433327 -0.01288564
## [5,]  0.09053546 -0.21476876 -0.61233625 -0.143458159  0.20354084  0.24132647
## [6,]  0.13536836  0.39338019 -0.13623518  0.518875575 -0.26320448  0.66744328
## [7,] -0.44266273  0.36461557 -0.02225060 -0.287362173  0.02320579  0.19215802
## [8,]  0.39567759  0.33073442  0.14046608 -0.357793846  0.14526955  0.18326804
## [9,]  0.58541387 -0.04395050  0.11321176 -0.006996245  0.08044702 -0.08325563
## [10,] -0.02289386  0.35524111 -0.10899245  0.543089950  0.63969694 -0.35621193
##          PC7          PC8          PC9          PC10
## [1,] -0.12505664  0.16000417 -0.09786284  0.0040664244
## [2,] -0.64536842 -0.50527345 -0.22592870 -0.0011724550
## [3,] -0.36922408  0.36326524  0.48502163 -0.1306454849
## [4,]  0.12583264 -0.44971705  0.52072259  0.0014025903
## [5,] -0.05986715  0.36382628 -0.56082145 -0.0003350351
## [6,]  0.14090663 -0.06320352  0.02286573 -0.0024913693
## [7,] -0.44342946 -0.01726032  0.01431296 -0.5945060184
## [8,] -0.28866482  0.46051936  0.33632123  0.3508546584
## [9,]  0.29575964  0.17583456  0.06581013 -0.7115949348
## [10,] -0.15757721  0.07233475 -0.01702224  0.0013460826
```

The feature with the strongest influence on the first PC is feature number 9 since this is the feature with the largest absolute value, and the feature with the weakest influence on the first PC is feature number 10 since this has the smallest absolute value.