

Week8_SimonsenHomework

Steven Simonsen

2024-10-15

```
library(dplyr)

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(caret)

## Loading required package: ggplot2
## Loading required package: lattice

library(NbClust)

setwd("C:\\Users\\steve\\OneDrive\\Documents\\School\\DSE6211\\Week8")

data <- read.csv("lab_8_data.csv")

training_ind <- createDataPartition(data$lodgepole_pine,
                                     p = 0.75,
                                     list = FALSE,
                                     times = 1)

training_set <- data[training_ind, ]
test_set <- data[-training_ind, ]

top_20_soil_types <- training_set %>%
  group_by(soil_type) %>%
  summarise(count = n()) %>%
  arrange(desc(count)) %>%
  select(soil_type) %>%
  top_n(20)

## Selecting by soil_type
training_set$soil_type <- ifelse(training_set$soil_type %in%
                                 top_20_soil_types$soil_type,
                                 training_set$soil_type,
                                 "other")
```

```

training_set$wilderness_area <- factor(training_set$wilderness_area)
training_set$soil_type <- factor(training_set$soil_type)

onehot_encoder <- dummyVars(~ wilderness_area + soil_type,
                             training_set[, c("wilderness_area",
                                                "soil_type")],
                             levelsOnly = TRUE,
                             fullRank = TRUE)

onehot_enc_training <- predict(onehot_encoder,
                               training_set[, c("wilderness_area",
                                                  "soil_type")])

training_set <- cbind(training_set, onehot_enc_training)

test_set$soil_type <- ifelse(test_set$soil_type %in%
                             top_20_soil_types$soil_type,
                             test_set$soil_type,
                             "other")

test_set$wilderness_area <- factor(test_set$wilderness_area)
test_set$soil_type <- factor(test_set$soil_type)

onehot_enc_test <- predict(onehot_encoder, test_set[,
                                                    c("wilderness_area",
                                                       "soil_type")])

test_set <- cbind(test_set, onehot_enc_test)

test_set[, -c(11:13)] <- scale(test_set[, -c(11:13)],
                              center = apply(training_set[,
                                                         -c(11:13)],
                                              2, mean),
                              scale = apply(training_set[,
                                                         -c(11:13)],
                                              2, sd))

training_set[, -c(11:13)] <- scale(training_set[, -c(11:13)])

training_features <- array(data = unlist(training_set[,
                                                    -c(11:13)]),
                           dim = c(nrow(training_set), 33))



training_labels <- array(data = unlist(training_set[, 13]),
                          dim = c(nrow(training_set)))

test_features <- array(data = unlist(test_set[, -c(11:13)]),
                       dim = c(nrow(test_set), 33))


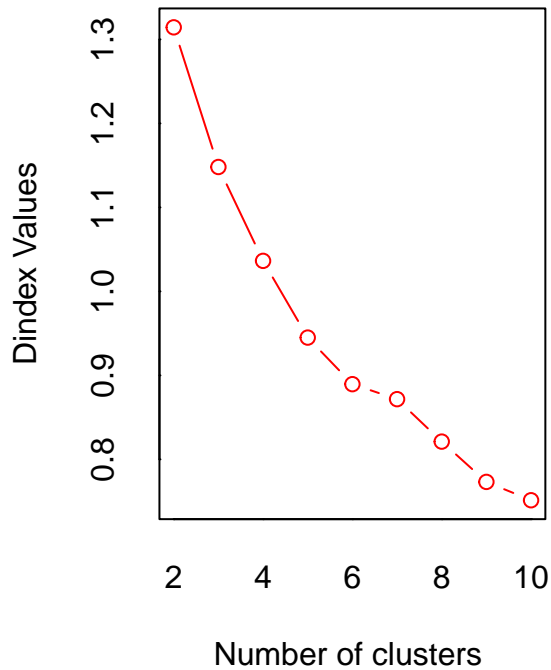
test_labels <- array(data = unlist(test_set[, 13]),
                      dim = c(nrow(test_set)))

```

```
set.seed(123)
nc <- NbClust(training_features[sample(nrow(training_features), 1000), c(4, 6, 10)],
              min.nc = 2, max.nc = 10, method = "kmeans")
```



```
## *** : The Hubert index is a graphical method of determining the number of clusters.
##       In the plot of Hubert index, we seek a significant knee that corresponds to a
##       significant increase of the value of the measure i.e the significant peak in Hubert
##       index second differences plot.
##
```



```
## *** : The D index is a graphical method of determining the number of clusters.
##           In the plot of D index, we seek a significant knee (the significant peak in Dindex
##           second differences plot) that corresponds to a significant increase of the value of
##           the measure.
```

```
## *****
## * Among all indices:
## * 6 proposed 2 as the best number of clusters
## * 5 proposed 3 as the best number of clusters
## * 7 proposed 4 as the best number of clusters
## * 4 proposed 9 as the best number of clusters
## * 1 proposed 10 as the best number of clusters
```

```
##           ***** Conclusion *****
```

```
## * According to the majority rule, the best number of clusters is 4
```

```
## *****
```

```
km_clusters <- kmeans(training_features[, c(4, 6, 10)], centers = 4)
```

```
cluster_number <- data.frame(cluster_number = km_clusters$cluster)
training_features <- cbind(training_features, cluster_number)
head(training_features)
```

```
##           1           2           3           4           5           6           7
```

```
## 1 -1.6114538 -0.5680384 -0.9792743 -0.41298603 -0.5009574 -1.0490066 0.7065639
## 2 -0.3640291 -0.2733056 -0.8437867 -0.00681951 -0.6894465 0.6082246 0.7822666
## 3 -1.2986989 1.5308164 -0.9792743 -1.13086174 -0.8093941 -1.1394245 -0.3154219
## 4 -0.1087924 -0.9967407 0.2401137 1.51866638 0.3044052 0.2271777 0.3659020
## 5 -0.1555259 0.6466180 -0.3018365 -0.95611568 -0.6894465 1.6977605 -0.3911246
## 6 -1.0542468 -0.8270460 -0.0308614 -0.27130004 -0.4838220 -0.7557942 0.7444152
##      8      9     10     11     12     13
## 1  0.3047499 -0.2846374 2.69809322 -0.2253618 -0.8939894 -0.2596715
## 2  0.6094299 -0.2064688 0.81638142 -0.2253618 -0.8939894 -0.2596715
## 3  0.5078699 0.6794428 -0.37077984 -0.2253618 -0.8939894 -0.2596715
## 4 -0.9647500 -0.7797057 -0.04262766 -0.2253618 -0.8939894 -0.2596715
## 5  1.4219098 1.1223986 -0.35165676 -0.2253618 -0.8939894 -0.2596715
## 6 -0.6092900 -0.8578743 -1.32999391 -0.2253618 -0.8939894 -0.2596715
##     14     15     16     17     18     19
## 1 -0.04795369 -0.03273853 -0.5001052 -0.08953925 -0.2349412 -0.2154418
## 2 -0.04795369 -0.03273853 -0.5001052 -0.08953925 -0.2349412 -0.2154418
## 3 -0.04795369 -0.03273853 -0.5001052 -0.08953925 -0.2349412 -0.2154418
## 4 -0.04795369 -0.03273853 1.9992736 -0.08953925 -0.2349412 -0.2154418
## 5 -0.04795369 -0.03273853 -0.5001052 -0.08953925 4.2557331 -0.2154418
## 6 -0.04795369 -0.03273853 -0.5001052 -0.08953925 -0.2349412 -0.2154418
##     20     21     22     23     24     25
## 1 -0.315538 -0.2933388 -0.05398665 -0.06195546 -0.01749279 -0.01236832
## 2 -0.315538 -0.2933388 -0.05398665 -0.06195546 -0.01749279 -0.01236832
## 3 -0.315538 -0.2933388 -0.05398665 -0.06195546 -0.01749279 -0.01236832
## 4 -0.315538 -0.2933388 -0.05398665 -0.06195546 -0.01749279 -0.01236832
## 5 -0.315538 -0.2933388 -0.05398665 -0.06195546 -0.01749279 -0.01236832
## 6 -0.315538 -0.2933388 -0.05398665 -0.06195546 -0.01749279 -0.01236832
##     26     27     28     29     30     31
## 1 -0.1604046 -0.1563453 -0.1390423 -0.1265158 -0.05105843 -0.1055234
## 2 -0.1604046 -0.1563453 -0.1390423 -0.1265158 -0.05105843 -0.1055234
## 3 -0.1604046 -0.1563453 -0.1390423 -0.1265158 -0.05105843 -0.1055234
## 4 -0.1604046 -0.1563453 -0.1390423 -0.1265158 -0.05105843 -0.1055234
## 5 -0.1604046 -0.1563453 -0.1390423 -0.1265158 -0.05105843 -0.1055234
## 6 -0.1604046 -0.1563453 -0.1390423 -0.1265158 -0.05105843 -0.1055234
##     32     33 cluster_number
## 1 -0.01749279 -0.04288122      4
## 2 -0.01749279 -0.04288122      3
## 3 -0.01749279 -0.04288122      2
## 4 -0.01749279 -0.04288122      1
## 5 -0.01749279 -0.04288122      3
## 6 -0.01749279 -0.04288122      2
```

Exercises

- 1) Why is it good practice to center and scale before applying k-means clustering? It is important to center using a mean of 0 and scale the data, typically to a standard deviation of 1, to avoid variables with larger magnitudes having more influence on the clusters, since they will dominate the distance calculation.
- 2) Print the cluster sizes and centers to the R console. Include a screenshot of the output.

```
km_clusters$size
```

```
## [1] 1268 3220 1326 723
```

```
km_clusters$centers
```

```
##           [,1]      [,2]      [,3]
## 1  1.47956864 -0.2889577 -0.02882224
## 2 -0.48366329 -0.6139885 -0.39437417
## 3 -0.21139514  1.2421734 -0.19289360
## 4 -0.05309446  0.9630976  2.16073076
```

3) Use the `aggregate()` function to calculate the mean of each variable within each cluster. Include a screenshot of the output.

```
aggregate(training_features[, -c(11:13)], by=list(training_features$cluster_number), mean)
```

```
##   Group.1      1      2      3      4      5
## 1      1  0.47307130 -0.017354926 -0.007460947  1.47956864  0.9190489
## 2      2 -0.39836362  0.009880297  0.176956314 -0.48366329 -0.2275790
## 3      3  0.48921083  0.084855563 -0.227553364 -0.21139514 -0.2189215
## 4      4  0.04727926 -0.169193621 -0.357680621 -0.05309446 -0.1967630
##           6      7      8      9     10     14
## 1 -0.2889577 -0.02019353  0.0167296  0.03247584 -0.02882224  0.083896263
## 2 -0.6139885 -0.05340142 -0.1246407 -0.05228389 -0.39437417 -0.002522781
## 3  1.2421734  0.02127219  0.2164019  0.13139752 -0.19289360 -0.047953688
## 4  0.9630976  0.23423379  0.1288811 -0.06508762  2.16073076 -0.047953688
##           15     16     17     18     19     20
## 1  0.13604048 -0.2418886 -0.07178512 -0.12161188  0.18670121  0.294479346
## 2 -0.03273853 -0.1927281  0.08524518 -0.06061378  0.03190053  0.004752669
## 3 -0.03273853  0.4442303 -0.08953925  0.16806803 -0.13853121 -0.144741782
## 4 -0.03273853  0.4678424 -0.08953925  0.17499587 -0.21544183 -0.272165701
##           21     22     23     24     25     26
## 1  0.278870672  0.03390412  0.0147012294  0.07268887 -0.01236832  0.04128903
## 2  0.009016867 -0.02514465 -0.0066134568 -0.01749279 -0.01236832 -0.07502304
## 3 -0.153751874  0.05807519 -0.0008690585 -0.01749279  0.04860582  0.15301146
## 4 -0.247257728 -0.05398665  0.0052649287 -0.01749279 -0.01236832 -0.01891219
##           27     28     29     30     31     32
## 1 -0.006508853 -0.07545402  0.266091643 -0.05105843 -0.05263340 -0.01749279
## 2 -0.040372072  0.10680851 -0.081630670  0.05259652  0.08787472 -0.01749279
## 3  0.105515815 -0.11140281 -0.059906465 -0.05105843 -0.10552340 -0.01749279
## 4 -0.002299684 -0.13904232  0.006753151 -0.05105843 -0.10552340  0.14066813
##           33 cluster_number
## 1 -0.04288122      1
## 2  0.04417298      2
## 3 -0.04288122      3
## 4 -0.04288122      4
```